

Solving Avalanche’s Multi-Coloring Problem via Ordering Constraint Graphs

John Boyd*

November 22, 2024

Abstract

Avalanche consensus achieves Byzantine agreement with constant-sized communication, enabling horizontal scaling to large validator sets. However, the protocol suffers from a multi-coloring problem: when conflict sets contain more than two choices, convergence probability decreases toward zero as choices increase. Existing solutions either sequentialize consensus or restrict transaction production, creating a trade-off between performance and decentralization. This work presents a novel solution by extracting pairwise ordering constraints from transactions to construct a constraint DAG, then executing Avalanche to reach agreement on the constraint graph rather than the transaction graph. By extracting only pairwise constraints — specifically, ordering relationships between pairs of parent transactions — we ensure each constraint conflicts with at most one other constraint (its opposite), guaranteeing multi-colored conflict sets cannot arise. Avalanche consensus therefore converges on the constraint DAG. Once peers agree on the ordering constraints, they deterministically derive total transaction ordering without additional communication. This approach preserves Avalanche’s $O(1)$ convergence time while maintaining full decentralization, at the cost of $O(n^2)$ communication overhead per transaction, where n is the number of parent transactions. Consequently, peers also reach agreement on total transaction ordering, a property not provided by Avalanche, enabling support for any general replicated state machine.

*john@coldnoise.net

1 Introduction

The Avalanche consensus protocol [3] is unique in its ability to reach Byzantine agreement with only constant-sized communication. This allows horizontal scaling without network partitioning—new validators may join to increase security and availability without degrading performance.

However, Avalanche suffers from a fundamental limitation known as the multi-coloring problem. When transactions form conflict sets with more than two options, the probability of convergence decreases as the number of options increases, potentially leading to liveness failures.

The Snowman consensus protocol addresses the multi-coloring problem through binary decomposition [1], but this requires executing multiple sequential instances of Avalanche consensus, hindering performance. The Snowman authors suggest the use of elected proposers to avoid multi-colored conflicts, reducing the probability of conflicts to mitigate this performance penalty.

This work proposes a novel solution to the multi-coloring problem that eliminates the need for sequential Avalanche instances while preserving the DAG structure and $O(1)$ convergence time, trading increased bandwidth for faster concurrent convergence. The key insight is to deliberately extract only pairwise ordering constraints from transactions to form a constraint DAG. By restricting constraints to pairs of transactions, we ensure each constraint $[a, b]$ conflicts only with $[b, a]$, making all conflict sets binary. Avalanche consensus is therefore guaranteed to converge on the constraint DAG. Once peers agree on the preferred constraints, they can deterministically compute the total transaction ordering.

An additional benefit of this construction is that peers achieve agreement on total transaction ordering, a property not provided by standard Avalanche. This enables the protocol to support general replicated state machines and other applications requiring transaction sequencing.

2 Background

2.1 Avalanche Consensus

Avalanche is a probabilistic, leaderless consensus protocol that uses network sampling to create a metastable mechanism driving participants to converge

on decisions [3]. The protocol records transactions into a DAG. As transactions disseminate, nodes insert them into the DAG and group them into conflict sets. Each node queries a randomly selected, constant-sized subset of peers for their preferred transaction in each conflict set. Preferred transactions receive chits and accumulate confidence scores. Once a transaction reaches the confidence threshold, it is considered final.

A node’s preference may change if subsequent queries indicate a transaction is no longer connected to the most preferred graph. This rule to adopt peer preferences makes indecision metastable and drives convergence.

While Avalanche provides agreement on which transactions are accepted, it does not provide agreement on the ordering of those transactions. Because transactions may be inserted simultaneously into the DAG, different peers may observe transactions in different orders, and Avalanche provides no mechanism to reconcile these differences.

2.2 The Multi-Coloring Problem

Avalanche consensus can be modeled as a graph coloring process over transaction conflict sets. Each conflict set is a vertex, and peers assign colors representing preferred transactions. Consensus is achieved when all peers agree on each vertex’s color.

The multi-coloring problem arises when conflict sets have more than two options. As described in [1], convergence probability decreases as the probability of two peers initially assigning the same color decreases. Assuming uniform random color assignment, convergence probability approaches zero as the number of colors increases.

For two-colored graphs, disagreement is metastable. The sampling process makes progress in expectation and eventually converges. This is not true for multi-colored graphs, posing a fundamental challenge to applications whose transactions do not form binary conflict sets.

2.3 Snowman’s Approach

Snowman consensus solves the multi-coloring problem through binary decomposition [1]. Each multi-colored conflict set of size L is decomposed into $\log_2(L)$ two-colored subsets. Snowman sequentially executes Avalanche consensus on each subset, and because each is binary, convergence is guaranteed.

The preferences form a path to exactly one transaction in the original conflict set.

While Snowman guarantees convergence, it increases convergence time by requiring multiple sequential Avalanche processes. To mitigate this, the authors suggest a proposer election process to reduce the probability of conflicts. Importantly, this conflicts with one of Avalanche’s defining properties: leaderless consensus. Concretely, this trade-off compromises decentralization to achieve performance.

3 Ordering Constraint Graphs

3.1 Core Idea

We solve the multi-coloring problem by extracting pairwise ordering constraints from transactions to form a constraint DAG. The critical design decision is to extract only binary constraints—specifically, ordering relationships between pairs of parent transactions. This is conceptually similar to the binary set decomposition of [1]; however, the use of ordering constraints allows us to create a constraint graph and make progress on each constraint concurrently, as opposed to sequentially. Intuitively, these partial ordering constraints directly relate to the underlying source of conflicts in the set (disagreement about order of observation).

This design ensures that each constraint $[p_0, p_1]$ conflicts only with its opposite $[p_1, p_0]$, making the maximum conflict set size exactly 2. Therefore, Avalanche consensus is guaranteed to converge on the constraint DAG. Since the constraint graph fully describes the transaction graph, by extension we also reach agreement on the transaction graph.

Trade-off: A transaction with n parents generates $\binom{n}{2} + n = O(n^2)$ constraints, each requiring Avalanche consensus. This increases communication overhead compared to standard Avalanche’s single consensus instance per transaction. However, these constraint instances execute concurrently rather than sequentially, preserving $O(1)$ convergence time in rounds while accepting higher bandwidth utilization. For applications where n is bounded by a small constant, this overhead is acceptable.

Multiple Avalanche instances must execute for a given transaction, but crucially they execute *concurrently* rather than sequentially. This allows convergence in the same time as standard Avalanche consensus, at the cost of

increased communication compared to basic Avalanche, but with significantly faster convergence than Snowman.

3.2 Notation

We use two DAGs: the *transaction DAG* whose vertices are transactions, and the *constraint DAG* whose vertices are ordering constraints. Transactions are annotated by height h and observation order i as $t_{h,i}$. For example, $t_{9,0}$ denotes the first transaction observed at height 9.

Ordering constraints are denoted $[t_a, t_b]$, where t_a is constrained to precede t_b . Constraints may be referenced as $c_{h,i,n}$, the n^{th} constraint of transaction $t_{h,i}$.

A special case exists for single-parent transactions: the *empty constraint* $[t_a, -]$. This does not actually constrain t_a 's ordering; it exists only to assist constraint DAG connectivity.

3.3 Constructing the Constraint DAG

Each transaction commits to a list of parent transactions which precede it, in the order in which they were observed by the node that produced this transaction. Ordering constraints are determined by transforming the transaction's list of parents into a list of ordered pairs. A transaction commits to:

1. The empty constraint for each parent
2. A pairwise ordering constraint for every ordered pair of parents consistent with the transaction's stated parent sequence

For example, transaction $t_{h,i}$ with parents p_0, p_1, p_2 (in that order) commits to: $[p_0, -]$, $[p_1, -]$, $[p_2, -]$, $[p_0, p_1]$, $[p_0, p_2]$, $[p_1, p_2]$. Each constraint implicitly depends on the parent transactions it references. These dependencies form the edges of the constraint DAG.

3.4 Conflict Sets in the Constraint DAG

To apply Avalanche consensus to the constraint DAG, we must define conflict sets. Two constraints $[a, b]$ and $[c, d]$ conflict if and only if $a = d$ and $b = c$. Because we extract only pairwise constraints, each constraint conflicts with




Figure 1: Example transaction DAG. Transactions $b_{2,0}$ and $b_{2,1}$ conflict due to their opposing parent orderings—they both reference $b_{1,0}$ and $b_{1,1}$ but in opposite order.

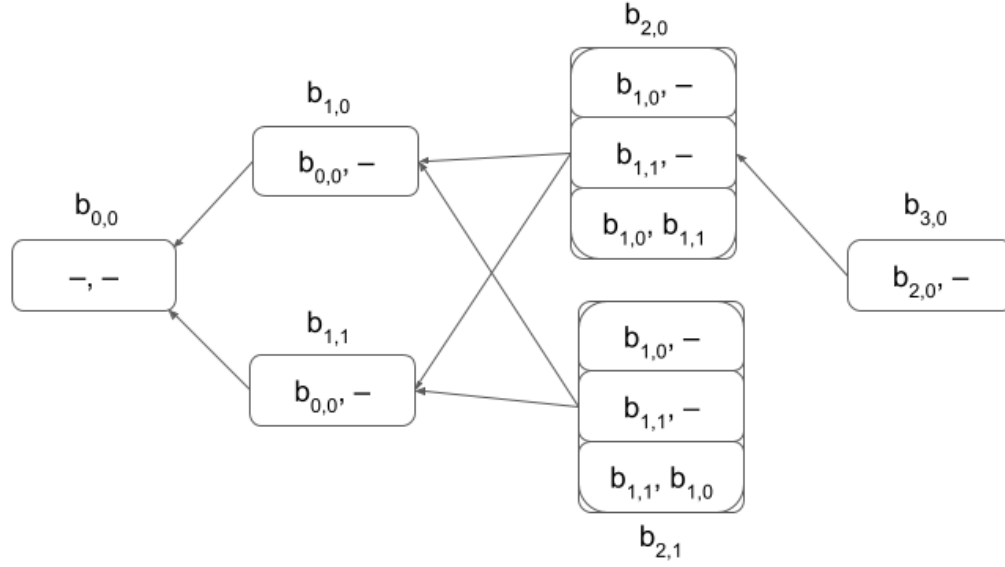


Figure 2: Expanded view showing parent ordering constraints for each transaction in Figure 1. Note that each constraint expresses an ordering relationship between exactly two parents.




Figure 3: Constraint DAG constructed from transactions in Figure 2. Conflict sets are indicated by dashed lines. Each conflict set contains at most two constraints.

at most one other constraint—its opposite. Constraints with no conflicts form singleton conflict sets.

Figure 3 illustrates the constraint DAG for our example. Constraints $[b_{1,0}, b_{1,1}]$ and $[b_{1,1}, b_{1,0}]$ are in the same conflict set because they reference parent transactions in opposing order. Avalanche consensus resolves the preferred constraint in each conflict set.

Consensus is reached by executing Avalanche on constraint conflict sets. Once every constraint committed by a transaction is accepted into the constraint DAG, that transaction is inferred to be accepted into the transaction DAG without any additional consensus process.

3.5 Handling Incompatible Constraints

An important property of this construction is that incompatible constraints can temporarily coexist in the constraint DAG during consensus. Consider the following scenario:

- Transaction A commits to constraint $[p_1, p_2]$

- Transaction B commits to constraint $[p_2, p_3]$
- Transaction C commits to constraint $[p_3, p_1]$

While these constraints are pairwise non-conflicting (none are direct opposites), they are mutually *incompatible*—they cannot all be satisfied in any valid total ordering. However, this does not create a structural problem for the protocol:

Any future transaction that attempts to build upon all three incompatible constraints will necessarily create a *direct conflict* with at least one of them. To extend the DAG beyond these three transactions, a new transaction must specify an ordering over some subset of $\{p_1, p_2, p_3\}$. Any such ordering will agree with at most two of the three incompatible constraints and will conflict with the third.

As new transactions are proposed and Avalanche consensus executes on their constraints, the metastability of indecision drives the network toward preferring transactions (and their constraints) that form a consistent ordering. The incompatible constraint that receives less support through this process will eventually be rejected, causing its parent transaction to be rejected as well.

This behavior is analogous to standard Avalanche: multiple conflicting transactions can temporarily exist in the DAG, but as the network builds upon preferred transactions, non-preferred transactions become disconnected from the preferred frontier and are eventually rejected.

Therefore, temporary coexistence of incompatible constraints does not threaten correctness or liveness. The protocol naturally resolves these incompatibilities through the standard Avalanche consensus mechanism.

3.6 Determining Transaction Sequence

As constraints are accepted, we can compute the final sequence of transactions. Once every constraint referencing any transaction in the progeny of transaction t is accepted, the sequence leading to t is computed recursively: merge the sequences at each parent (eliminating duplicates), then append each parent in order, as shown in Algorithm 4.

When two transactions disagree on parent ordering, they produce conflicting pairwise constraints. Eventually one constraint is rejected from the constraint DAG, at which point its transaction is rejected from the transac-


```

1: procedure SEQUENCEAT( $t$ )
2:    $\mathcal{S} := \emptyset$ 
3:   for  $t' \in \mathcal{T} : t' \xleftarrow{*} t$  do
4:     for  $t'' \in \mathcal{T} : t'' \in \text{SEQUENCEAT}(t') \wedge t'' \notin \mathcal{S}$  do
5:        $\mathcal{S} = \mathcal{S} || t''$ 
6:   for  $t' \in \mathcal{T} : t' \xleftarrow{*} t$  do
7:      $\mathcal{S} = \mathcal{S} || t'$ 
8:   return  $\mathcal{S}$ 

```

Figure 4: Algorithm to determine the sequence of transactions preceding a given transaction.

tion DAG. This ensures the transaction graph can never disagree with the constraint graph.

Applying this algorithm to Figure 1 to determine the sequence at $b_{3,0}$, we obtain: $b_{0,0}$, $b_{1,0}$, $b_{1,1}$, $b_{2,0}$, $b_{3,0}$.

4 Analysis

Table 1 provides an overview of the comparison of this work with Avalanche consensus and Snowman.

Property	This Work	Avalanche [3]	Snowman [1]
Liveness	Guaranteed	Fails with multi-coloring	Guaranteed
Time to Convergence	$O(1)$	$O(1)$ (no multi-color)	$O(\log L)$
Communication Complexity	$O(n^2)$ per transaction	$O(1)$ per transaction	$O(\log L)$ transactions
Decentralization	Full	Full	Reduced (proposers)

Table 1: Comparison of protocol characteristics. L is the size of a multi-colored conflict set, n is the number of parent transactions per transaction.

4.1 Correctness

The key property ensuring correctness is that our construction guarantees all conflict sets in the constraint DAG are binary. This is achieved by deliberately extracting only pairwise ordering constraints from each transaction. Since a pairwise constraint $[a, b]$ can only conflict with its opposite $[b, a]$, each conflict set contains at most two elements.

Since Avalanche consensus is proven to converge for binary conflict sets [3], and our constraint DAG contains only binary conflict sets by construction, the overall process is guaranteed to converge.

Once the constraint DAG reaches consensus, the transaction sequence is deterministically derivable by all peers using Algorithm 4. Therefore, all peers will agree on the total ordering of transactions.

4.2 Liveness

This construction strengthens the liveness guarantees of Avalanche. Under the assumed threshold of honest participants, Avalanche guarantees liveness when the transaction graph only forms binary conflict sets. Liveness is not guaranteed, however, for transaction graphs that may form non-binary conflict sets (i.e., the multi-coloring problem). Our construction eliminates the possibility of non-binary conflict sets, thereby guaranteeing liveness under the same assumptions.

4.3 Safety

This construction inherits the safety properties of Avalanche consensus [3]. Specifically, the protocol is secure against Byzantine faults for any adversarial presence $f' \leq \frac{n(k-\alpha-\Psi)}{k} \leq f$, where f' is the number of adversarial peers and n, k, α, Ψ are protocol parameters.

4.4 Performance

Time to Convergence: This work maintains the same convergence time as standard Avalanche (assuming no multi-colored conflicts). While multiple Avalanche instances must execute (one per constraint), they run concurrently rather than sequentially. In contrast, Snowman requires $\log_2(L)$ sequential rounds of consensus for a conflict set of size L , significantly increasing latency.

Communication Complexity: Each transaction with n parents generates $O(n^2)$ pairwise ordering constraints: specifically, n empty constraints and $\binom{n}{2} = \frac{n(n-1)}{2}$ pairwise ordering constraints. Each constraint requires Avalanche consensus, which involves querying a constant-sized sample of k peers per round for $O(1)$ expected rounds. Therefore, the communication complexity per transaction is $O(n^2)$, where n is the number of parents. In contrast, standard Avalanche requires $O(1)$ communication per transaction (assuming no multi-colored conflicts), and Snowman requires $O(\log L)$ sequential consensus instances for a conflict set of size L .

5 Applications

The primary contribution of this work is solving Avalanche’s multi-coloring problem while preserving its performance characteristics and decentralization. However, an important consequence of the constraint-based approach is that peers achieve agreement on total transaction ordering.

Standard Avalanche provides agreement on which transactions are accepted but not on their ordering. This limits its applicability to scenarios where transaction order does not matter or where conflicts can be resolved through transaction execution. This is well suited for UTXO-style transaction ledgers, as in Bitcoin [2], but not suitable for smart-contract execution protocols that require strict transaction sequencing.

By achieving total ordering of transactions, this work is now suitable for smart-contract execution platforms, and indeed any application that conforms to the replicated state machine model.

6 Conclusion

This work presents a novel solution to Avalanche’s multi-coloring problem through the use of ordering constraint graphs. The key insight is to deliberately extract only pairwise ordering constraints from transactions, ensuring that all conflict sets in the resulting constraint DAG are binary. This guarantees convergence under Avalanche consensus without compromising performance or decentralization.

This approach offers significant advantages over existing solutions: it maintains Avalanche’s $O(1)$ convergence time by executing constraint con-

sensus concurrently, avoids the need to restrain transaction production, and maintains full decentralization.

As a consequence of this construction, peers also achieve agreement on total transaction ordering, a property not provided by standard Avalanche. This enables the protocol to support general replicated state machines and other applications requiring transaction sequencing, significantly expanding the applicability of Avalanche consensus.

The technique may be applied to improve existing Avalanche-based protocols or enable new applications requiring both horizontal scalability and transaction ordering guarantees.

References

- [1] Aaron Buchwald, Stephen Buttolph, Andrew Lewis-Pye, Patrick O’Grady, and Kevin Sekniqi. Frosty: Bringing strong liveness guarantees to the snow family of consensus protocols, 2024.
- [2] Satoshi Nakamoto. Bitcoin: A peer-to-peer electronic cash system, 2008.
- [3] Team Rocket, Maofan Yin, Kevin Sekniqi, Robbert van Renesse, and Emin Gün Sirer. Scalable and probabilistic leaderless bft consensus through metastability, 2020.